

课程名称：操作系统实践	年级：2023级	上机实践成绩：
指导教师：张民	姓名：李彤	
上机实践名称：忙等待实验	学号：10235101500	上机实践日期：2024.11.04
上机实践编号：	组号：	上机实践时间：

展示忙等待

原因分析

当前线程主动将 CPU 执行权释放，然后经过一段时间（ticks）后，系统将该线程唤醒，将其重新加入到就绪队列中等待调度。但是由于调度队列是优先级调度，导致在一定时间内，该线程被反复加入就绪队列（ready）以及从队列中取出执行（running），这样的过程非常占用 CPU 并且唤醒顺序非常混乱。所以线程状态从running变成 ready 会造成 CPU 忙等待。

观察发现的一些问题

在运行最开始的代码时，可以发现，所有`time_ticks()`的返回值都是负数，显然不符合逻辑，所以我在`yield()`中的调试语句中使用强制类型转换将其值转化成非负数，并且将后续`check_and_wakeup_sleep_thread`中的`time_ticks()`返回值也进行强制转化。

```
Pintos hda1
Loading.....
Kernel command line: -q run alarm-multiple
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... Yield:thread main at tick -4611106404000858108.
Yield:thread main at tick -4611106404000858104.
Yield:thread main at tick -4611106404000858100.
Yield:thread main at tick -4611106404000858096.
Yield:thread main at tick -4611106404000858092.
130,867,200 loops/s.
Boot complete.
Executing 'alarm-multiple':
(alarm-multiple)Yield:thread main at tick -4611106404000858088.
  begin
(alarm-multiple) Creating 5 threads to sleep 7 times each.
(alarm-multiple) Thread 0 sleeps 10 ticks each time,
(alarm-multiple) thread 1 sleeps 20 ticks each time, and so on.
(alarm-multiple) If successful, product of iteration count and
(alarm-multiple) sleep duration will appear in nondescending order.
Yield:thread main at tick -4611106404000858088.
Yield:thread thread 0 at tick -4611106404000858088.
Yield:thread thread 1 at tick -4611106404000858087.
Yield:thread thread 2 at tick -4611106404000858087.
Yield:thread thread 3 at tick -4611106404000858087.
Yield:thread thread 4 at tick -4611106404000858087.
Yield:thread main at tick -4611106404000858087.
Yield:thread thread 0 at tick -4611106404000858087.
Yield:thread thread 1 at tick -4611106404000858087.
Yield:thread thread 2 at tick -4611106404000858087.
```

修改代码

```

void thread_yield(void)
{
    struct thread *cur = thread_current();
    enum intr_level old_level;

    ASSERT(!intr_context());

    old_level = intr_disable();
    if (cur != idle_thread)
    {
        // list_push_back (&ready_list, &cur->elem);
        list_insert_ordered(&ready_list, &cur->elem, prio_cmp_func, NULL);
        printf("Yield:thread %s at tick %lld.\n", cur->name, (uint64_t)timer_ticks());
        cur->status = THREAD_READY;
        schedule();
        intr_set_level(old_level);
    }
}

```

这是修正后的结果

```

Pintos hda1
Loading.....
Kernel command line: -q run alarm-multiple
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... Yield:thread main at tick 4.
Yield:thread main at tick 8.
Yield:thread main at tick 12.
Yield:thread main at tick 16.
Yield:thread main at tick 20.
Yield:thread main at tick 24.
116,121,600 loops/s.
Boot complete.
Executing 'alarm-multiple':
(alarm-multiple) begin
(alarm-multiple) Creating 5 threads to sleep 7 times each.
(alarm-multiple) Thread 0 sleeps 10 ticks each time,
(alarm-multiple) thread 1 sleeps 20 ticks each time, and so on.
(alarm-multiple) If successful, product of iteration count and
(alarm-multiple) sleep duration will appear in nondescending order.
Yield:thread main at tick 27.
Yield:thread thread 0 at tick 27.
Yield:thread thread 1 at tick 27.
Yield:thread thread 2 at tick 27.
Yield:thread thread 3 at tick 27.
Yield:thread thread 4 at tick 27.
Yield:thread main at tick 27.
Yield:thread thread 0 at tick 27.
Yield:thread thread 1 at tick 27.
Yield:thread thread 2 at tick 27.

```

通过观察忙等待的输出可以发现，线程的调度并没有按照优先级实现。这里先放一下，等到最后和正确代码比较。

```

(alarm-multiple) thread 4: duration=50, iteration=2, product=100
(alarm-multiple) thread 1: duration=20, iteration=5, product=100

```

实验步骤

1. 首先添加两个我们需要用到的变量

其中`thread_status`用来记录每一个线程的状态，在这里我们添加`THREAD_SLEEP`用来表示进程处于 休眠状态

。

在`thread`结构体中添加变量`wake_time`用来记录线程 **休眠结束** 的时间。

```
enum thread_status
{
    THREAD_RUNNING,    /**< Running thread. */
    THREAD_READY,      /**< Not running but ready to run. */
    THREAD_BLOCKED,    /**< Waiting for an event to trigger. */
    THREAD_DYING,      /**< About to be destroyed. */
    THREAD_SLEEP       /**< Sleppng thread,litong*/
};

int64_t wake_time;    /**< Time when thread wake up. litong*/
};
```

2. 进程休眠实现

先修改`timer_sleep`函数的内部接口

```
void
timer_sleep (int64_t ticks)
{
    // int64_t start = timer_ticks ();

    // ASSERT (intr_get_level () == INTR_ON);
    // while (timer_elapsed (start) < ticks)
    //     thread_yield ();

    thread_sleep(ticks);
}
```

接下来，我们就要实现`thread_sleep`函数了。

`thread_sleep`函数首先判断了休眠时长是否合法，如果不合法，就直接退出。然后获取当前进程，`cur != idle_thread`确保了CPU不是在空等待，接着设置进程的状态和休眠结束的时间，`schedule`函数调度进程，将进程插入 **ready队列** 而不是直接执行，也不是插入waiting队列。等到下次调用`ready`队列中的进程时，再按照

优先级重新调度

```
void thread_sleep(int64_t ticks){
    if(ticks <= 0) return;
    struct thread *cur = thread_current();

    // 禁用中断并保存当前中断级别
    enum intr_level old_level = intr_disable();

    if (cur != idle_thread)    // 确保CPU不是在空等待
    {
        cur->status = THREAD_SLEEP;        // 将当前进程状态改为休眠
        cur->wake_time = timer_ticks() + ticks;    // 设置进程休眠结束的时间
        schedule();        // 调度进程，将进程插入ready队列而不是直接执行，也不是插入waiting队列
    }

    // 恢复之前的中断级别
    intr_set_level(old_level);
}
```

3. 唤醒进程

还是先介绍一下函数`check_and_wakeup_sleep_thread()`，这个函数会遍历当前所有进程，并判断是否有进程休眠结束，如果有，那么它就会把这个进程按照优先级有序地插入ready队列等待执行，同时输出一些调试信息。

```
void check_and_wakeup_sleep_thread(){
    struct list_elem *e = list_begin(&all_list);
    int64_t cur_ticks = (uint64_t)timer_ticks();

    // 遍历所有线程列表
    while(e != list_end(&all_list)){
        struct thread *t = list_entry(e, struct thread, allelem);
        enum intr_level old_level = intr_disable();

        // 如果线程处于睡眠状态且当前时间已达到或超过唤醒时间
        if(t->status == THREAD_SLEEP && cur_ticks >= t->wake_time) {
            t->status = THREAD_READY;

            // 将线程插入到就绪队列中，按优先级排序
            list_insert_ordered(&ready_list, &t->elem, prio_cmp_func, NULL);
            printf("Wake up thread %s at tick %lld.\n", t->name, cur_ticks); // 提示信息
        }

        // 移动到下一个进程
        e = list_next(e);
        // 恢复之前的中断级别
        intr_set_level(old_level);
    }
}
```

因此，我们可以通过把函数`check_and_wakeup_sleep_thread`加到时钟中断里面，这样每发生一次时钟中断，我们就可以排查一次进程链表，确保那些休眠结束的进程能够及时进入等待序列（不是waiting，仍指

ready)

```
/** Timer interrupt handler. */
static void
timer_interrupt (struct intr_frame *args UNUSED)
{
    ticks++;
    thread_tick ();

    // 检查当前时间是否到了线程的睡眠时间，如果是，则唤醒该线程
    check_and_wakeup_sleep_thread();
}
```

抢占式优先级调度的实现

这里我主要是在`check_and_wakeup_sleep_thread`函数内进行修改。

主要思路是，用一个thread指针，记录被唤醒线程中最高优先级线程，等到遍历唤醒结束以后，再将该线程的优先级与当前running线程的优先级进行比较，如果该优先级更大，就发生抢占，否则不抢占。具体实现代码如下：

```
void check_and_wakeup_sleep_thread(void) {
    int64_t cur_ticks = timer_ticks(); // 获取当前滴答数
    struct thread *highest_priority_thread = NULL; // 记录被唤醒线程中最高优先级线程

    // 遍历所有线程列表
    struct list_elem *e = list_begin(&all_list);
    while (e != list_end(&all_list)) {
        enum intr_level old_level = intr_disable(); // 在整个操作中禁用中断，减少频繁的中断开关

        struct thread *t = list_entry(e, struct thread, allelem);
        e = list_next(e); // 提前获取下一个线程，防止后续对当前线程状态的修改影响遍历

        // 如果线程处于休眠状态且当前时间已达到或超过唤醒时间
        if (t->status == THREAD_SLEEP && cur_ticks >= t->wake_time) {
            // 禁用中断以进行安全操作
            enum intr_level old_level = intr_disable();

            t->status = THREAD_READY; // 将线程状态设置为就绪
            list_insert_ordered(&ready_list, &t->elem, prio_cmp_func, NULL); // 插入到就绪队列

            // 更新被唤醒线程中最高优先级线程
            if (!highest_priority_thread || t->priority > highest_priority_thread->priority) {
                highest_priority_thread = t;
            }
            printf("Wake up thread %s at tick %lld.\n", t->name, cur_ticks); // 输出日志

            // 恢复中断状态
            intr_set_level(old_level);
        }
    }
}
```

```

    }
}

// 抢占逻辑
if (highest_priority_thread) {
    struct thread *current = thread_current(); // 获取当前运行线程
    // 如果唤醒的线程优先级高于当前线程，进行抢占
    if (highest_priority_thread->priority > current->priority) {
        thread_yield();
    }
}
}
}

```

一开始我的抢占函数直接用的是`thread_yield()`函数，但是发现会爆如下的错误，大概就是上下文切换的问题，在网上查询资料后才知道这里还需要确保线程不在中断上下文中，因此我们可以用`intr_yield_on_return`函数替换，它会在中断返回时立即触发调度，切换到优先级更高的线程，保证了调度的安全。

（也可以在`thread_yield()`函数前加一个`if(!intr_context())`的判断语句，确保线程不在中断上下文中。这种方法输出除了少了一条 `Yield:thread idle is at ...` 语句其他都差不多，所以我觉得这两种方法应该都可选

```

Loading.....
Kernel command line: -q run alarm-multiple
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... Kernel PANIC at ../../threads/thread.c:344 in thread_yield(): assertion '!intr_context()' failed.
Call stack: 0xc00296f8 0xc0020efb 0xc0020af0 0xc0024013 0xc0021b7f 0xc0021d9f 0xc0024031 0xc0023dc6 0xc002033e.
The 'backtrace' program can make call stacks useful.
Read "Backtraces" in the "Debugging Tools" chapter
of the Pintos documentation for more information.
Timer: 1 ticks
Thread: 0 idle ticks, 1 kernel ticks, 0 user ticks

```

```

// 抢占逻辑
if (highest_priority_thread) {
    struct thread *current = thread_current(); // 获取当前运行线程
    // 如果唤醒的线程优先级高于当前线程，进行抢占
    if (highest_priority_thread->priority > current->priority) {
        // thread_yield();
        intr_yield_on_return();
    }
}
}

```

但是似乎还是有点问题，`running_thread`竟然都是idle，cpu在空等待，令人费解。

```
(alarm-multiple) sleep duration will appear in nondescending order.
Wake up thread thread 0 at tick 134.
Yield:thread idle at tick 134.
Wake up thread thread 0 at tick 144.
Wake up thread thread 1 at tick 144.
Yield:thread idle at tick 144.
Wake up thread thread 0 at tick 154.
Wake up thread thread 2 at tick 154.
Yield:thread idle at tick 154.
Wake up thread thread 0 at tick 164.
Wake up thread thread 1 at tick 164.
Wake up thread thread 3 at tick 164.
Yield:thread idle at tick 164.
```

在经过7200s的调试后 疑似发现问题出在那里了，我给代码加了一段调试语句：

```
printf("Ready list contents:\n");
for (e = list_begin(&ready_list); e != list_end(&ready_list); e = list_next(e)) {
    struct thread *t = list_entry(e, struct thread, elem);
    printf("Thread %s with priority %d\n", t->name, t->priority);
}
```

输出结果如下：

```
Ready list contents:
Ready list contents:
Ready list contents:
Ready list contents:
Ready list contents:
Ready list contents:
Wake up thread thread 3 at tick 286.
```

ready_list竟然是空的！！！！这下就能解释为什么前面cpu都在空转了——压根没有进程在等待执行，发现问题之后，就要准备解决问题了。

按照同样的方法，我也输出了一下all_list链表，似乎也有点不对劲：

```
All list contents:
Thread with priority 0
Thread with priority 0
Thread with priority 0
Ready list contents:
All list contents:
Thread with priority 0
Thread with priority 0
Thread with priority 0
Ready list contents:
```

and 在原本的输出日志里多加了priority的输出，发现所有进程的优先级都相同（所以抢占不了）

```
printf("Wake up thread %s at tick %lld with priority %d.\n", t->name, cur_ticks, t->priority); // 输出日志
```



```
Wake up thread thread 3 at tick 363 with priority 31.  
Yield:thread idle at tick 363.  
Wake up thread thread 4 at tick 373 with priority 31.  
Yield:thread idle at tick 373.  
Wake up thread thread 3 at tick 403 with priority 31.
```

由此我们可以知道，如果要实现抢占式的优先级调度，那必然要先设计算法使创建的thread优先级能够发生改变。

测试代码

before

前面讲到在运行原始代码时，进程并没有按照正确的顺序被调度，这是因为忙等待导致进程被唤醒顺序非常混乱。

```
(alarm-multiple) thread 4: duration=50, iteration=2, product=100  
(alarm-multiple) thread 1: duration=20, iteration=5, product=100
```

after

在修改代码之后，系统能够按照正确的优先级顺序调度进程。

```
(alarm-multiple) thread 3: duration=40, iteration=2, product=80  
(alarm-multiple) thread 2: duration=30, iteration=3, product=90  
(alarm-multiple) thread 1: duration=20, iteration=5, product=100  
(alarm-multiple) thread 4: duration=50, iteration=2, product=100  
(alarm-multiple) thread 1: duration=20, iteration=6, product=120  
(alarm-multiple) thread 2: duration=30, iteration=4, product=120
```

为了确保实验结果的正确性，我还跑了多次测试，结果都符合预期。下面给出了其中两次测试的结果。


```

Pintos hda1
Loading.....
Kernel command line: -q run alarm-multiple
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... Yield:thread main at tick 4.
Yield:thread main at tick 8.
Yield:thread main at tick 12.
Yield:thread main at tick 16.
Yield:thread main at tick 20.
104,755,200 loops/s.
Boot complete.
Executing 'alarm-multiple':
(alarm-multiple) begin
(alarm-multiple) Creating 5 threads to sleep 7 times each.
(alarm-multiple) Thread 0 sleeps 10 ticks each time,
(alarm-multiple) thread 1 sleeps 20 ticks each time, and so on.
(alarm-multiple) If successful, product of iteration count and
(alarm-multiple) sleep duration will appear in nondescending order.
Wake up thread thread 0 at tick 130.
Wake up thread thread 0 at tick 140.
Wake up thread thread 1 at tick 140.
Wake up thread thread 0 at tick 150.
Wake up thread thread 2 at tick 150.
Wake up thread thread 0 at tick 160.
Wake up thread thread 1 at tick 160.
Wake up thread thread 3 at tick 160.
Wake up thread thread 0 at tick 170.
Wake up thread thread 4 at tick 170.
Wake up thread thread 0 at tick 180.
Wake up thread thread 1 at tick 180.

```

```

Pintos hda1
Loading.....
Kernel command line: -q run alarm-multiple
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... Yield:thread main at tick 4.
Yield:thread main at tick 8.
Yield:thread main at tick 12.
Yield:thread main at tick 16.
Yield:thread main at tick 20.
Yield:thread main at tick 24.
114,073,600 loops/s.
Boot complete.
Executing 'alarm-multiple':
(alarm-multiple) begin
(alarm-multiple) Creating 5 threads to sleep 7 times each.
(alarm-multiple) Thread 0 sleeps 10 ticks each time,
(alarm-multiple) thread 1 sleeps 20 ticks each time, and so on.
(alarm-multiple) If successful, product of iteration count and
(alarm-multiple) sleep duration will appear in nondescending order.
Yield:thread main at tick 28.
Wake up thread thread 0 at tick 137.
Wake up thread thread 0 at tick 147.
Wake up thread thread 1 at tick 147.
Wake up thread thread 0 at tick 157.
Wake up thread thread 2 at tick 157.
Wake up thread thread 0 at tick 167.
Wake up thread thread 1 at tick 167.
Wake up thread thread 3 at tick 167.
Wake up thread thread 0 at tick 177.
Wake up thread thread 4 at tick 177.

```

遗留问题

1. 输出乱序

但是，我发现在最终的输出中，总是会有一些很“碍眼”的存在，如下

```
(alarm-multiple) thread 1: duration=20, iteration=7, product=140
(alarm-multiple) thread 2: duration=30, iteration=5, product=150
(alarm-multiple) thread 4: duration=5Yield:thread main at tick 581.
0, iteration=3, product=150
(alarm-multiple) thread 5: duration=40, iteration=4, product=180
(alarm-multiple) thread 2: duration=30, iteration=6, product=180
(alarm-multiple) thread 3: duration=40, iteration=5, product=200
```

可以看出，在这一段的输出中发生了混乱，两条输出语句同时进行，这就导致了内容交错输出。

结合我之前的了解（看过一点xv6的代码），我认为是某个语句块没有加锁，导致在输出时CPU被其他进程抢占，这就造成了内容交替输出的局面，后面还去问了助教（and热心同学的解答），确实是多线程的问题。

这里先保留问题，课后还可以再了解一下多线程和锁的相关内容。

2. 线程优先级的修改和ready_list的插入

因为近期考试比较多，所以没花太多时间继续编写、调试关于线程优先级的修改和ready_list的插入的代码。

关于抢占式优先级调度的实现基本就是前面代码所示，最终的结果大概就是下面这样，如果能实现线程优先级的修改，应该是能够实现抢占了。

```
Wake up thread thread 2 at tick 304 with priority 31.
Wake up thread thread 3 at tick 324 with priority 31.
Wake up thread thread 4 at tick 324 with priority 31.
Wake up thread thread 2 at tick 334 with priority 31.
Wake up thread thread 3 at tick 364 with priority 31.
Wake up thread thread 4 at tick 374 with priority 31.
Wake up thread thread 3 at tick 404 with priority 31.
Wake up thread thread 4 at tick 424 with priority 31.
Wake up thread thread 4 at tick 474 with priority 31.
Wake up thread main at tick 574 with priority 31.
(alarm-multiple) thread 0: duration=10, iteration=1, product=10
(alarm-multiple) thread 0: duration=10, iteration=2, product=20
(alarm-multiple) thread 1: duration=20, iteration=1, product=20
(alarm-multiple) thread 0: duration=10, iteration=3, product=30
(alarm-multiple) thread 2: duration=30, iteration=1, product=30
(alarm-multiple) thread 0: duration=10, iteration=4, product=40
(alarm-multiple) thread 1: duration=20, iteration=2, product=40
(alarm-multiple) thread 3: duration=40, iteration=1, product=40
(alarm-multiple) thread 0: duration=10, iteration=5, product=50
(alarm-multiple) thread 4: duration=50, iteration=1, product=50
(alarm-multiple) thread 0: duration=10, iteration=6, product=60
(alarm-multiple) thread 1: duration=20, iteration=3, product=60
```

这里ready_list为什么会是空链表我也暂时还没有找到原因

后续又调试了一下，发现ready_list不是空链表，还是有插入的，只不过比较少，所以都被其他调试语句淹没了。那现在剩下的问题应该就只有线程优先级的修改了。

```
if(list_empty(&ready_list))  
    printf("Empty\n");  
else printf("Something\n");
```

```
Empty  
Wake up thread thread 4 at tick 475 with priority 31.  
Something  
Empty
```

心得体会

通过这个实验，我主要加深了对下面这些事物的了解和认识：

- 进程休眠的实现原理，以及休眠结束后进程的去向，进程调度的大致流程；
- 出现忙等待的原因；
- `thread_yield()`和`thread_sleep()`两个函数的底层实现以及功能差异；
- 抢占式优先级调度的实现原理，以及发生抢占时所需要的条件。

虽然每次都是看着课件做的实验，但都多多少少会出现一些问题，随着对操作系统理解的深入，现在已经能够自己根据出现的Error调试修改代码了，而不是像之前那样什么都先抛到浏览器上。