

课程名称：操作系统实践	年级：2023级	上机实践成绩：
指导教师：张民	姓名：李彤	
上机实践名称：忙等待实验	学号：10235101500	上机实践日期：2024.11.04
上机实践编号：	组号：	上机实践时间：

## 展示忙等待

### 原因分析

当前线程主动将 CPU 执行权释放，然后经过一段时间（ticks）后，系统将该线程唤醒，将其重新加入到就绪队列中等待调度。但是由于调度队列是优先级调度，导致在一定时间内，该线程被反复加入就绪队列（ready）以及从队列中取出执行（running），这样的过程非常占用 CPU 并且唤醒顺序非常混乱。所以线程状态从running变成 ready 会造成 CPU 忙等待。

### 观察发现的一些问题

在运行最开始的代码时，可以发现，所有`time_ticks()`的返回值都是负数，显然不符合逻辑，所以我在`yield()`中的调试语句中使用强制类型转换将其值转化成非负数，并且将后续`check_and_wakeup_sleep_thread`中的`time_ticks()`返回值也进行强制转化。

```
Pintos hda1
Loading.....
Kernel command line: -q run alarm-multiple
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... Yield:thread main at tick -4611106404000858108.
Yield:thread main at tick -4611106404000858104.
Yield:thread main at tick -4611106404000858100.
Yield:thread main at tick -4611106404000858096.
Yield:thread main at tick -4611106404000858092.
130,867,200 loops/s.
Boot complete.
Executing 'alarm-multiple':
(alarm-multiple)Yield:thread main at tick -4611106404000858088.
  begin
(alarm-multiple) Creating 5 threads to sleep 7 times each.
(alarm-multiple) Thread 0 sleeps 10 ticks each time,
(alarm-multiple) thread 1 sleeps 20 ticks each time, and so on.
(alarm-multiple) If successful, product of iteration count and
(alarm-multiple) sleep duration will appear in nondescending order.
Yield:thread main at tick -4611106404000858088.
Yield:thread thread 0 at tick -4611106404000858088.
Yield:thread thread 1 at tick -4611106404000858087.
Yield:thread thread 2 at tick -4611106404000858087.
Yield:thread thread 3 at tick -4611106404000858087.
Yield:thread thread 4 at tick -4611106404000858087.
Yield:thread main at tick -4611106404000858087.
Yield:thread thread 0 at tick -4611106404000858087.
Yield:thread thread 1 at tick -4611106404000858087.
Yield:thread thread 2 at tick -4611106404000858087.
```

### 修改代码

```

void thread_yield(void)
{
    struct thread *cur = thread_current();
    enum intr_level old_level;

    ASSERT(!intr_context());

    old_level = intr_disable();
    if (cur != idle_thread)
    {
        // list_push_back (&ready_list, &cur->elem);
        list_insert_ordered(&ready_list, &cur->elem, prio_cmp_func, NULL);
        printf("Yield:thread %s at tick %lld.\n", cur->name, (uint64_t)timer_ticks());
        cur->status = THREAD_READY;
        schedule();
        intr_set_level(old_level);
    }
}

```

这是修正后的结果

```

Pintos hda1
Loading.....
Kernel command line: -q run alarm-multiple
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... Yield:thread main at tick 4.
Yield:thread main at tick 8.
Yield:thread main at tick 12.
Yield:thread main at tick 16.
Yield:thread main at tick 20.
Yield:thread main at tick 24.
116,121,600 loops/s.
Boot complete.
Executing 'alarm-multiple':
(alarm-multiple) begin
(alarm-multiple) Creating 5 threads to sleep 7 times each.
(alarm-multiple) Thread 0 sleeps 10 ticks each time,
(alarm-multiple) thread 1 sleeps 20 ticks each time, and so on.
(alarm-multiple) If successful, product of iteration count and
(alarm-multiple) sleep duration will appear in nondescending order.
Yield:thread main at tick 27.
Yield:thread thread 0 at tick 27.
Yield:thread thread 1 at tick 27.
Yield:thread thread 2 at tick 27.
Yield:thread thread 3 at tick 27.
Yield:thread thread 4 at tick 27.
Yield:thread main at tick 27.
Yield:thread thread 0 at tick 27.
Yield:thread thread 1 at tick 27.
Yield:thread thread 2 at tick 27.

```

通过观察忙等待的输出可以发现，线程的调度并没有按照优先级实现。这里先放一下，等到最后和正确代码比较。

```

(alarm-multiple) thread 4: duration=50, iteration=2, product=100
(alarm-multiple) thread 1: duration=20, iteration=5, product=100

```

## 实验步骤

### 1. 首先添加两个我们需要用到的变量

其中`thread_status`用来记录每一个线程的状态，在这里我们添加`THREAD_SLEEP`用来表示进程处于休眠状态。

在`thread`结构体中添加变量`wake_time`用来记录线程 **休眠结束** 的时间。

```
enum thread_status
{
    THREAD_RUNNING,    /**< Running thread. */
    THREAD_READY,      /**< Not running but ready to run. */
    THREAD_BLOCKED,    /**< Waiting for an event to trigger. */
    THREAD_DYING,      /**< About to be destroyed. */
    THREAD_SLEEP       /**< Sleppng thread,litong*/
};

int64_t wake_time;    /**< Time when thread wake up. litong*/
};
```

## 2. 进程休眠实现

先修改`timer_sleep`函数的内部接口

```
void
timer_sleep (int64_t ticks)
{
    // int64_t start = timer_ticks ();

    // ASSERT (intr_get_level () == INTR_ON);
    // while (timer_elapsed (start) < ticks)
    //     thread_yield ();

    thread_sleep(ticks);
}
```

接下来，我们就要实现`thread_sleep`函数了。

`thread_sleep`函数首先判断了休眠时长是否合法，如果不合法，就直接退出。然后获取当前进程，`cur != idle_thread`确保了CPU不是在空等待，接着设置进程的状态和休眠结束的时间，`schedule`函数调度进程，将进程插入 **ready队列** 而不是直接执行，也不是插入waiting队列。等到下次调用ready队列中的进程时，再按照

## 优先级重新调度

```
void thread_sleep(int64_t ticks){
    if(ticks <= 0) return;
    struct thread *cur = thread_current();

    // 禁用中断并保存当前中断级别
    enum intr_level old_level = intr_disable();

    if (cur != idle_thread)    // 确保CPU不是在空等待
    {
        cur->status = THREAD_SLEEP;        // 将当前进程状态改为休眠
        cur->wake_time = timer_ticks() + ticks;    // 设置进程休眠结束的时间
        schedule();        // 调度进程，将进程插入ready队列而不是直接执行，也不是插入waiting队列
    }

    // 恢复之前的中断级别
    intr_set_level(old_level);
}
```

## 3. 唤醒进程

还是先介绍一下函数`check_and_wakeup_sleep_thread()`，这个函数会遍历当前所有进程，并判断是否有进程休眠结束，如果有，那么它就会把这个进程按照优先级有序地插入ready队列等待执行，同时输出一些调试信息。

```
void check_and_wakeup_sleep_thread(){
    struct list_elem *e = list_begin(&all_list);
    int64_t cur_ticks = (uint64_t)timer_ticks();

    // 遍历所有线程列表
    while(e != list_end(&all_list)){
        struct thread *t = list_entry(e, struct thread, allelem);
        enum intr_level old_level = intr_disable();

        // 如果线程处于睡眠状态且当前时间已达到或超过唤醒时间
        if(t->status == THREAD_SLEEP && cur_ticks >= t->wake_time) {
            t->status = THREAD_READY;

            // 将线程插入到就绪队列中，按优先级排序
            list_insert_ordered(&ready_list, &t->elem, prio_cmp_func, NULL);
            printf("Wake up thread %s at tick %lld.\n", t->name, cur_ticks); // 提示信息
        }

        // 移动到下一个进程
        e = list_next(e);
        // 恢复之前的中断级别
        intr_set_level(old_level);
    }
}
```

因此，我们可以通过把函数`check_and_wakeup_sleep_thread`加到时钟中断里面，这样每发生一次时钟中断，我们就可以排查一次进程列表，确保那些休眠结束的进程能够及时进入等待序列（不是waiting，仍指

Ready)

```
/** Timer interrupt handler. */
static void
timer_interrupt (struct intr_frame *args UNUSED)
{
    ticks++;
    thread_tick ();

    // 检查当前时间是否到了线程的睡眠时间，如果是，则唤醒该线程
    check_and_wakeup_sleep_thread();
}
```

## 测试代码

### before

前面讲到在运行原始代码时，进程并没有按照正确的顺序被调度，这是因为忙等待导致进程被唤醒顺序非常混乱。

```
(alarm-multiple) thread 4: duration=50, iteration=2, product=100
(alarm-multiple) thread 1: duration=20, iteration=5, product=100
```

### after

在修改代码之后，系统能够按照正确的优先级顺序调度进程。

```
(alarm-multiple) thread 5: duration=40, iteration=2, product=80
(alarm-multiple) thread 2: duration=30, iteration=3, product=90
(alarm-multiple) thread 1: duration=20, iteration=5, product=100
(alarm-multiple) thread 4: duration=50, iteration=2, product=100
(alarm-multiple) thread 1: duration=20, iteration=6, product=120
(alarm-multiple) thread 2: duration=30, iteration=4, product=120
```

为了确保实验结果的正确性，我还跑了多次测试，结果都符合预期。下面给出了其中两次测试的结果。

```

Pintos hda1
Loading.....
Kernel command line: -q run alarm-multiple
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... Yield:thread main at tick 4.
Yield:thread main at tick 8.
Yield:thread main at tick 12.
Yield:thread main at tick 16.
Yield:thread main at tick 20.
104,755,200 loops/s.
Boot complete.
Executing 'alarm-multiple':
(alarm-multiple) begin
(alarm-multiple) Creating 5 threads to sleep 7 times each.
(alarm-multiple) Thread 0 sleeps 10 ticks each time,
(alarm-multiple) thread 1 sleeps 20 ticks each time, and so on.
(alarm-multiple) If successful, product of iteration count and
(alarm-multiple) sleep duration will appear in nondescending order.
Wake up thread thread 0 at tick 130.
Wake up thread thread 0 at tick 140.
Wake up thread thread 1 at tick 140.
Wake up thread thread 0 at tick 150.
Wake up thread thread 2 at tick 150.
Wake up thread thread 0 at tick 160.
Wake up thread thread 1 at tick 160.
Wake up thread thread 3 at tick 160.
Wake up thread thread 0 at tick 170.
Wake up thread thread 4 at tick 170.
Wake up thread thread 0 at tick 180.
Wake up thread thread 1 at tick 180.

```

```

Pintos hda1
Loading.....
Kernel command line: -q run alarm-multiple
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... Yield:thread main at tick 4.
Yield:thread main at tick 8.
Yield:thread main at tick 12.
Yield:thread main at tick 16.
Yield:thread main at tick 20.
Yield:thread main at tick 24.
114,073,600 loops/s.
Boot complete.
Executing 'alarm-multiple':
(alarm-multiple) begin
(alarm-multiple) Creating 5 threads to sleep 7 times each.
(alarm-multiple) Thread 0 sleeps 10 ticks each time,
(alarm-multiple) thread 1 sleeps 20 ticks each time, and so on.
(alarm-multiple) If successful, product of iteration count and
(alarm-multiple) sleep duration will appear in nondescending order.
Yield:thread main at tick 28.
Wake up thread thread 0 at tick 137.
Wake up thread thread 0 at tick 147.
Wake up thread thread 1 at tick 147.
Wake up thread thread 0 at tick 157.
Wake up thread thread 2 at tick 157.
Wake up thread thread 0 at tick 167.
Wake up thread thread 1 at tick 167.
Wake up thread thread 3 at tick 167.
Wake up thread thread 0 at tick 177.
Wake up thread thread 4 at tick 177.

```

## 遗留问题

但是，我发现在最终的输出中，总是会有一些很“碍眼”的存在，如下

```
(alarm-multiple) thread 1: duration=20, iteration=7, product=140
(alarm-multiple) thread 2: duration=30, iteration=5, product=150
(alarm-multiple) thread 4: duration=5Yield:thread main at tick 581.
0, iteration=3, product=150
(alarm-multiple) thread 5: duration=40, iteration=4, product=100
(alarm-multiple) thread 2: duration=30, iteration=6, product=180
(alarm-multiple) thread 3: duration=40, iteration=5, product=200
```

可以看出，在这一段的输出中发生了混乱，两条输出语句同时进行，这就导致了内容交错输出。

结合我之前的了解（看过一点xv6的代码），我认为是某个语句块没有加锁，导致在输出时CPU被其他进程抢占，这就造成了内容交替输出的局面，后面还去问了助教（and热心同学的解答），确实是多线程的问题。

这里先保留问题，课后还可以再了解一下多线程和锁的相关内容。

## 心得体会

通过这个实验，我主要加深了对下面这些事物的了解和认识：

- 进程休眠的实现原理，以及休眠结束后进程的去向，进程调度的大致流程；
- 出现忙等待的原因；
- `thread_yield()`和`thread_sleep()`两个函数的底层实现以及功能差异。

感觉现在做实验越来越熟练了o(￣▽￣)ブ。虽然每次都是看着课件做的实验，但都多多少少会出现一些问题，随着对操作系统理解的深入，现在已经能够自己根据出现的Error调试修改代码了，而不是像之前那样什么都先抛到浏览器上。